




1

About this course	
Name	Vận hành hệ thống học máy (Machine Learning Operations)
Course ID	IT5414
Credit	2
Slides, group registration, and announcement	– On Microsoft Teams – y8izumb
Time & location	

MLOps is a combination of data science and DevOps principles that speeds up model deployment and ensures product quality. This course provides an overview of the MLOps process, including planning, developing, testing, deploying, managing, and operating machine learning models in a production environment


ĐẠI HỌC BÁCH KHOA HÀ NỘI
 HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

2

Evaluation

- Project: (including coding, report, representation)
 - Teamwork : 2-3 students/group
 - Topics : to be defined
- Final examination



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

3

3

ML research vs. ML production

	Research	Production
Objectives	Model performance*	Different stakeholders have different objectives



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY


*** It's actively being worked. See [Utility is in the Eye of the User: A Critique of NLP Leaderboards](#) (Ethayarajh and Jurafsky, EMNLP 2020)

4

4

Stakeholder objectives

ML team
highest




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

5


5

Stakeholder objectives

ML team
highest



Sales
sells more




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

6


6

Stakeholder objectives

ML team
highest



Sales
sells more



Product
fastest




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

7


7

Stakeholder objectives

ML team
highest



Sales
sells more



Product
fastest



Manager
maximizes profit



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

8

8

Leaderboard-style ML

- More comprehensive utility function
 - Model performance (e.g. accuracy)
 - Latency
 - Prediction cost
 - Interpretability
 - Robustness
 - Ease of use (e.g. OSS tools, community support)
 - Hardware requirements
- Adaptive to different use cases
 - Instead of a leaderboard for each dataset/task, the leaderboard adapts to each company's needs
- Dynamic datasets
 - Realistic distribution shifts with different types of shifts



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

9

9

Computational priority

	Research	Production
Objectives	Model performance	Different stakeholders have different objectives
Computational priority	Fast training, high throughput	Fast inference , low latency

generating predictions



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY


10

10

Latency matters


- 100ms delay can hurt conversion rates by 7% ([Akamai study](#) '17)
- 30% increase in latency costs 0.5% conversion rate ([Booking.com](#) '19)
- 53% phone users will leave a page that takes >3s to load ([Google](#) '16)

11



- Latency: time to move a leaf
- Throughput: how many leaves in 1 sec

12



13

ML in research vs. in production

	Research	Production
Objectives	Model performance	Different stakeholders have different objectives
Computational priority	Fast training, high throughput	Fast inference, low latency
Data	Static	Constantly shifting

14

Data

Research	Production
• Clean • Static • Mostly historical data	• Messy • Constantly shifting • Historical + streaming data • Biased, and you don't know how biased • Privacy + regulatory concerns

THE COGNITIVE CODER

By **Armand Ruiz**, Contributor, InfoWorld | SEP 26, 2017 7:22 AM PDT


The 80/20 data science dilemma

Most data scientists spend only 20 percent of their time on actual data analysis and 80 percent of their time finding, cleaning, and reorganizing huge amounts of data, which is an inefficient data strategy

ML in research vs. in production

	Research	Production
Objectives	Model performance	Different stakeholders have different objectives
Computational priority	Fast training, high throughput	Fast inference, low latency
Data	Static	Constantly shifting
Fairness	Good to have (sadly)	Important

Fairness



Google Shows Men Ads for Better Jobs

by Krista Bradford | Last updated Dec 1, 2019



The Berkeley study found that both face-to-face and online lenders rejected a total of 1.3 million creditworthy black and Latino applicants between 2008 and 2015. Researchers said they believe the applicants "would have been accepted had the applicant not been in these minority groups." That's because when they used the income and credit scores of the rejected applications but deleted the race identifiers, the mortgage application was accepted.

ML in research vs. in production

	Research	Production
Objectives	Model performance	Different stakeholders have different objectives
Computational priority	Fast training, high throughput	Fast inference, low latency
Data	Static	Constantly shifting
Fairness	Good to have (sadly)	Important
Interpretability*	Good to have	Important



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

19

19

Interpretability



Zoom poll: which one would you want as your surgeon?




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

20

20

Interpretability



Result from the Zoom poll
last year

1. Who would you rather pick?

AI Surgeon (90% accuracy) (44) 67%

Human Surgeon (80% accuracy) (22) 33%



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

21

21

ML in research vs. in production

	Research	Production
Objectives	Model performance	Different stakeholders have different objectives
Computational priority	Fast training, high throughput	Fast inference, low latency
Data	Static	Constantly shifting
Fairness	Good to have (sadly)	Important
Interpretability	Good to have	Important



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY


22

22

Traditional software

Separation of Concerns is a design principle for separating a computer program into distinct sections such that each section addresses a separate concern

- Code and data are separate
 - Inputs into the system shouldn't change the underlying code



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
Image by [Arda Cetinkaya](#)

23

23

ML systems

- Code and data are tightly coupled
 - ML systems are part code, part data
- Not only test and version code, need to test and version data too

the hard
part

ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY


24

24

Test and version data

- Extremely hard to ensure correctness in time

Don't panic.
We'll revisit
later!



ĐẠI HỌC BÁCH KHOA HÀ NỘI

HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

[Back to the Future: Solving the time-travel problem in machine learning - Tecton](#)

25

25

ML systems: version data

- Line-by-line diffs like Git doesn't work with datasets
- Can't naively create multiple copies of large datasets
- How to merge changes?



ĐẠI HỌC BÁCH KHOA HÀ NỘI

HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

26

26

How to ...

- Validate data correctness?
- Test features' usefulness?
- Detect when the underlying data distribution has changed?
- Know if the changes are bad for models without ground truth labels?
- Detect malicious data?
 - Not all data points are equal (e.g. scans of cancerous lungs are more valuable)
 - Bad data might harm your model and/or make it susceptible to attacks

ML systems: data poisoning attacks




Fig. 1: An illustrating example of backdoor attacks. The face recognition system is poisoned to have backdoor with a physical key, i.e., a pair of commodity reading glasses. Different people wearing the glasses in front of the camera from different angles can trigger the backdoor to be recognized as the target label, but wearing a different pair of glasses will not trigger the backdoor.



REMEMBER WHEN A BILLION-PARAMETER MODEL WAS CONSIDERED LARGE

PEPPERIDGE FARM REMEMBERS

SWITCH TRANSFORMERS: SCALING TO TRILLION PARAMETER MODELS WITH SIMPLE AND EFFICIENT SPARSITY

 **ĐẠI HỌC BÁCH KHOA HÀ NỘI**
 HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

William Fedus*
 Google Brain
 liamfedus@google.com

Barret Zoph*
 Google Brain
 barretzoph@google.com

Noam Shazeer
 Google Brain
 noam@google.com

29

29

Engineering challenges with large ML models

- Too big to fit on-device
- Consume too much energy to work on-device
- Too slow to be useful
 - Autocompletion is useless if it takes longer to make a prediction than to type
- If unit/CI tests take hours, the development cycles will stagnate

 **ĐẠI HỌC BÁCH KHOA HÀ NỘI**
 HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

30

30

5. ML production myths



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

31

31

Myth #1: Deploying is hard



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

32

32

Myth #1: Deploying is hard

Deploying is easy. Deploying reliably is hard



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

33

33

Myth #2: You only deploy one or two ML models at a time




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

34

34

Myth #2: You only deploy one or two ML models at a time

Booking.com: 150+ models, Uber: thousands




Myth #3: You won't need to update your models as much

DevOps: Pace of software delivery is accelerating


- Elite performers deploy **973x** more frequently with **6570x** faster lead time to deploy ([Google DevOps Report, 2021](#))
- DevOps standard (2015)
 - Etsy deployed 50 times/day
 - Netflix 1000s times/day
 - AWS every 11.7 seconds

DevOps to MLOps: Slow vs. Fast

Only 11% of organizations can put a model into production within a week, and 64% take a month or longer



We'll learn how to do minute-iteration cycle!



Accelerating ML Delivery



How
often **SHOULD**
I update
my models?



How often
CAN I update
my models?



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

39

39

ML + DevOps = 



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

40

40

Myth #4: ML can magically transform your business overnight



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

41

41

Myth #4: ML can magically transform your business overnight

Magically: possible
Overnight: no




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

42

42

Efficiency improves with maturity

Model deployment timeline and ML maturity



[2020 state of enterprise machine learning](#) (Algorithmia, 2020)

ML engineering is more engineering than ML

MLEs might spend most of their time:

- wrangling data
 - understanding data
 - setting up infrastructure
 - deploying models
- instead of training ML models



Myth #5: Most ML engineers don't need to worry about scale




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

45

45

Myth #5: Most ML engineers don't need to worry about scale

Company Size



71,791 responses

[StackOverflow Developer Survey 2019](#)

ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

46

46

What this course is not

- How different machine learning models works




Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

47

47

Let's start where it all began

Machine learning in production is fantastic
BUT




Massive technical depth is incurred if not careful

Hidden Technical Debt in Machine Learning Systems

D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips
{dsulley, gholit, dgg, edavydov, toddphilips}@google.com
Google, Inc.

Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, Dan Dennison
{ebner, vchaudhary, mwyong, jfcrespo, dennison}@google.com
Google, Inc.




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

48

48

Why do we focus on modelling?




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

49

49

Machine Learning in the real world




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

50


50

Key observations

1. Machine Learning in production is much more than doing ML modelling



2. Machine Learning in production is a cycle



51

The other stuff is DevOps

DevOps = Developer operations

- Dates to late 80s and early 90s
- Around 2007/2008 rose to popularity to remove the separation between software development with its operations part/IT department

💡 This is both a joke and not.

💡 MLOps is directly derived from DevOps.

💡 Therefore, let's try to understand DevOps first.



52

So, what is DevOps?

This is the closest to a definition that I could find:

DevOps is a set of practices that combines software development (*Dev*) and IT operations (*Ops*). It aims to shorten the systems development life cycle and provide continuous delivery with high software quality. It's an combination of human mindset, processes and technologies that continuously creates value.



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY


53

53

So, what is DevOps?

This is the closest to a definition that I could find:

DevOps is a set of practices that combines software development (*Dev*) and IT operations (*Ops*). It aims to shorten the systems development **life cycle** and provide continuous delivery with high software quality. It's an combination of **human mindset**, **processes** and **technologies** that continuously creates value.




Technologies
ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

54

54

Technology, Processes, Mindset




Use technologies that support the different parts of the lifecycle

Implement processes to make sure everyone is in sync about the lifecycle

Always consider all part of the lifecycle, not just its parts

But then MLOps must be...

Is a set of **tools**, **processes**, and **mindset** that aim to make **ML Lifecycle** reproducible, trackable, testable and maintainable to continuously create value.




Let's look at the different phases of MLOps

Data phase

- 🔥 Business understanding
- 🔥 Data understanding
- 🔥 Designing the ML-powered software

Machine Learning Operations (MLOps)




57

Model phase

- 🔥 Model engineering
- 🔥 Data engineering
- 🔥 Deliver a stable quality ML model that we will run in production

Machine Learning Operations (MLOps)




58

Operations phase

- 🔥 Deliver the previously developed ML model in production
- 🔥 Testing, versioning, continuous delivery, and monitoring

Machine Learning Operations (MLOps)




59

If DevOps exist, then why do we need MLOps?



Because data changes everything




60

What makes an MLOps engineer?

A mix of

- Software developing
- Machine Learning
- Data engineering




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

61


61

Where's waldo?



Size of bubble =
communication /
technical writing

Where should the different
positions be?



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

62

62

According to stable diffusion

<https://stablediffusionweb.com/>

Prompt:

"Machine Learning operations engineer"


A MLOps engineer is

🔥 Buff

🔥 Locked in

🔥 Many screens

Its not completely wrong



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

63


63

Why does companies care about MLOps

Having automated model deployed with errors can cost A LOT of money:

"A famous example of the dangers here was Knight Capital's system losing \$465 millions in 45 minutes, apparently because of unexpected behavior from obsolete experimental codepaths"

– Hidden Technical depth in Machine Learning Systems



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

64

64

Why is MLOps hot?

1. Growing Adoption of Machine Learning (ML)
2. Complexity of ML Workflow
3. Bridge between Development and Operations
4. Need for Collaboration
5. Ensuring Model Governance and Compliance
6. Automation and Scalability
7. Continuous Integration and Continuous Deployment (CI/CD)
8. Infrastructure Orchestration
9. Adoption of Cloud Services
10. Business Impact



ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

65

65

Open AI study

🔍 What is the contributes to the success of OpenAI?

- 💡 Funding
- 💡 Data
- 💡 People
- 💡 Compute
- 💡 Service contract with Microsoft

Microsoft fired 10,000 workers in the same breath as the invested \$10B in OpenAI




ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

66

66

Trends in MLOps

MLOps has been trending for a couple of
Tools have been the main priority




67

Choosing the right tool for the job

68

Looking back

MLOps around 2006 = write everything from scratch



Today we have options


- There is a tool for everything you need

Pros:

- + Easy to get started
- + Easy to iterate


Cons:

- Framework integration can be really hard
- Hard to compare frameworks



MLOps now


Pick a *stack* of tools



71

Not picking the right tool leads to TD


- In a nutshell MLOps is about dealing reducing technical debt
 - ⚠️ Technical debt is the implied cost of future reworking required when choosing an easy but limited solution instead of a better approach that could take more time



72


MLOps is full stack

- In MLOps we embrace the full stack of problems that comes from the full lifecycle. Especially integration problems.
- Criteria for what goes into the stack (4Cs):
 - 💡 Cost
 - 💡 Coverage
 - 💡 Complexity
 - 💡 Community
- Whenever we need to pick one tool over the other, we need to consider these 4 criteria.
- And most time this is not possible without actually trying to use both.



MLOps at its core is...

- ...delivering value for business 🏢
- ...thinking about the hole pipeline, not just data and model 🤖
- ...accounting for long term goals from the start 📅



Meme of the day



75

References

1. Machine Learning Engineering with Python by Andrew P. McMahan
2. Implementing MLOps in the Enterprise by Yaron Haviv and Noah Gift
3. Machine Learning Engineering in Action by Ben Wilson
4. Designing Machine Learning Systems by Chip Huyen
5. Practical MLOps by Noah Gift and Alfredo Deza
6. Reliable Machine Learning by Cathy Chen et al.
7. Treveil, Mark, et al. Introducing MLOps. O'Reilly Media, 2020.
8. Sculley, David, et al. "Hidden technical debt in machine learning systems." Advances in neural information processing systems 28 (2015).
9. MLOps: Continuous delivery and automation pipelines in machine learning.
<https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>

76