



Multimodal Machine Learning

Lecture 5.1: Multimodal Transformers (Part 1)


Dam Quang Tuan

Objectives of today's class


- Positional embeddings
- Language pre-training
 - BERT: Bidirectional Encoder Representations from Transformers
- Multimodal transformers (Image and language)
 - Concatenated transformers (VisualBERT, Uniter)
 - Crossmodal transformers (ViLBERT, LXMERT)
 - Modality-shift transformer (MAG-BERT)
- Sequence-to-sequence modeling with Transformers

Positional Embeddings


Transformer Self-Attention



Transformer Multi-Head Self-Attention



Transformer Multi-Head Self-Attention




What happens if the words are shuffled?

Position embeddings

❑ Position information is not encoded in a self-attention module

How can we encode position information?

Simple approach: one-hot encoding




Position embeddings


- ❑ Position information is not encoded in a self-attention module

How can we encode position information?

Simple approach: one-hot encoding + linear embeddings + $\left\{ \begin{array}{l} \text{Sum} \\ \text{- or -} \\ \text{concat} \end{array} \right.$




Transformer Multi-Head Self-Attention




Transformer Multi-Head Self-Attention


In vector format...



Transformer Multi-Head Attention




Transformer – Residual Connection




Language Pre-training

Token-level and Sentence-level Embeddings


Token-level embeddings



Sentence-level embedding




Pre-Training and Fine-Tuning



BERT: Bidirectional Encoder Representations from Transformers

Advantages:

- ① Jointly learn representation for token-level and sentence level
- ② Same network architecture for pre-training and fine-tuning




BERT: Bidirectional Encoder Representations from Transformers

Advantages:

- ① Jointly learn representation for token-level and sentence level
- ② Same network architecture for pre-training and fine-tuning
- ③ Can be used learn relationship between sentences
- ④ Models bidirectional and long-range interactions between tokens


How can
we do all
this?



BERT: Bidirectional Encoder Representations from Transformers

Advantages:

- ① Jointly learn representation for token-level and sentence level
- ② Same network architecture for pre-training and fine-tuning
- ③ Can be used learn relationship between sentences
- ④ Models bidirectional interactions between tokens




Pre-training BERT Model

1 Masked Language Model

Randomly mask input tokens and then try to predict them

What is the loss function?



Pre-training BERT Model


② Next Sentence Prediction

Given two sentences, predict if this is the next one or not

What is the loss function?

Where can we find training data?

How can BERT know the difference between both sentences?



Fine-Tuning BERT

4 Question-answering: find start/end of the answer in the document


Paragraph: “... Other legislation followed, including the Migratory Bird Conservation Act of 1929, a 1937 treaty prohibiting the hunting of right and gray whales, and the Bald Eagle Protection Act of 1940. These later laws had a low cost to society—the species were relatively rare—and little opposition was raised.”

Question 1: “Which laws faced significant *opposition*?”

Plausible Answer: *later laws*

Question 2: “What was the name of the 1937 *treaty*?”


Plausible Answer: *Bald Eagle Protection Act*




How?

Fine-Tuning BERT

- 4 Question-answering: find start/end of the answer in the document




Other Fine-tuning Approaches




Language-Vision Transformers

Multimodal Embeddings



Option 1: Concatenate modalities and learn BERT transformer


Simple Solution: Contextualized Multimodal Embeddings



VisualBERT




A person hits a ball with a tennis racket




Li, Liunian Harold, et al. "Visualbert: A simple and performant baseline for vision and language." *arXiv* (2019).

UNITER

Similar Transformer architecture to BERT and VisualBERT... but with slightly different optimization




Multimodal Embeddings




Option 2: Look at pairwise interactions between modalities


Multimodal Transformer – Pairwise Cross-Modal




Cross-Modal Transformer Module ($V \rightarrow L$)




Cross-Modal Transformer Module ($V \rightarrow L$)




Cross-Modal Transformer Module ($\beta \rightarrow \alpha$)



ViLBERT



LXMERT



Tan, Hao, and Mohit Bansal. "Lxmert: Learning cross-modality encoder representations from transformers." *arXiv* (August 20, 2019).

Reminder: Modality-Shifting Fusion





Figure 1: Conceptual figure demonstrating that the word representation of the same underlying word “sick” can vary conditioned on different co-occurring nonverbal behaviors. The nonverbal context during a word segment is depicted by the sequence of facial expressions and intonations. The speaker who has a relatively soft voice and frowning behaviors at the second time step displays negative sentiment.

Reminder: Modality-Shifting Fusion



Example with language modality:


Primary modality: language

Secondary modalities: acoustic and visual

Negative-shifted representation

word: "expectations"

Positive-shifted representation



Reminder: Modality-Shifting Fusion





Figure 2: An illustrative example for Recurrent Attended Variation Embedding Network (RAVEN) model: The RAVEN model has three components: (1) Nonverbal Sub-networks, (2) Gated Modality-mixing Network, and (3) Multimodal Shifting. For the given word “sick” in the utterance, the Nonverbal Sub-networks first computes the visual and acoustic embedding through modeling the sequence of visual and acoustic features lying in a word-long segment with separate LSTM network. The Gated Modality-mixing Network module then infers the nonverbal shift vector as the weighted average over the visual and acoustic embedding based on the original word embedding. The Multimodal Shifting finally generates the multimodal-shifted word representation by integrating the nonverbal shift vector to the original word embedding. The multimodal-shifted word representation can be then used in the high-level hierarchy to predict sentiments or emotions expressed in the sentence.


Modality-Shifting with Transformers

Multimodal Adaptation Gate (MAG) + BERT





Sequence-to-Sequence Using Transformer




Sequence-to-Sequence Modeling



Seq2Seq with Transformer Attentions





Seq2Seq with Transformer Attentions



Seq2Seq with Transformer Attentions

How should we connect the encoder and decoder self-attention to the transformer attention?



Seq2Seq with Transformer Attentions

